

Introduction to Machine Learning

AST 5765

Michael Himes

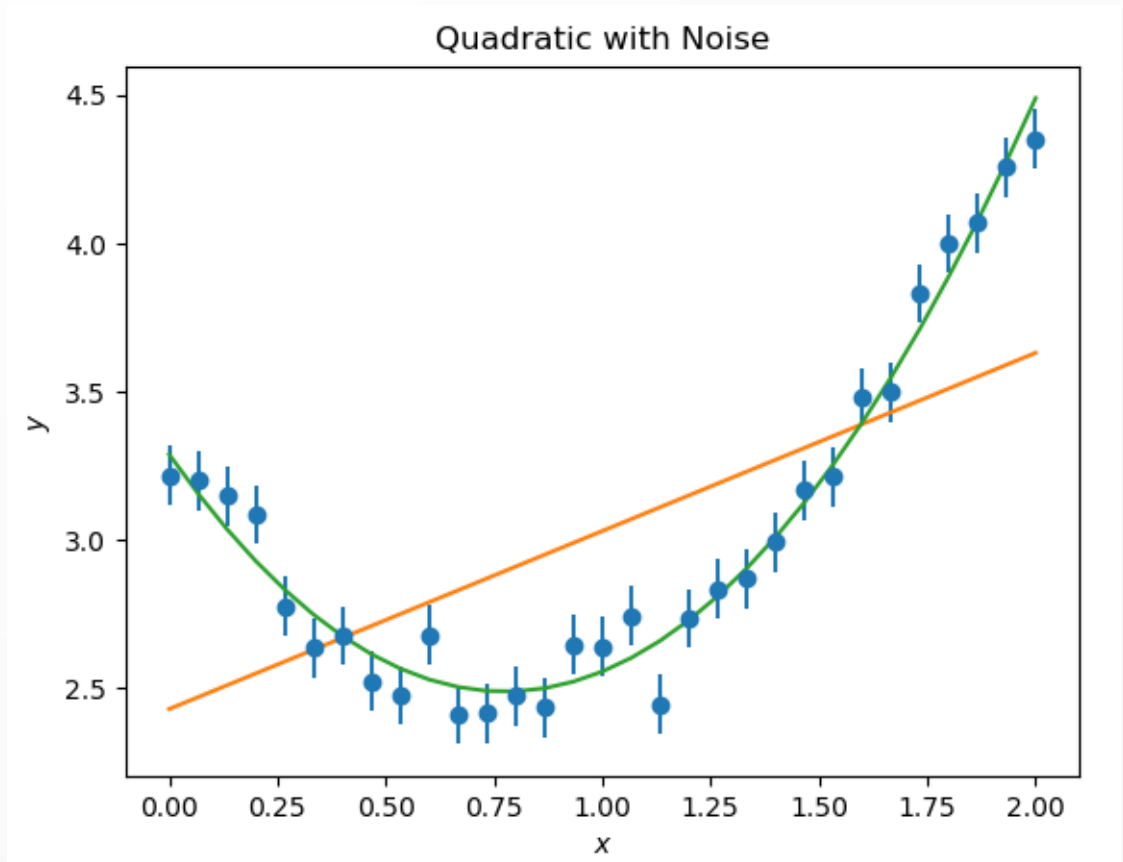
November 19, 2020

Overview

- Review model fitting
- What is machine learning?
- Types of ML
- ML details
- Demo: visual neural network
- ML applications
- Demo: basic dense neural networks


Model fitting

- If we have a data set of (x, y) pairs, we can fit the parameters of some model
- Least squares
 - $y_i = p_1 x_i + p_0$
 - $y_i = p_2 x_i^2 + p_1 x_i + p_0$
- Modeling a complicated function?
 - $y_i = \sum_n p_n x_i^n$




Model fitting

- What if we have a data set of inputs ($x_{1,i}$, $x_{2,i}$, ..., $x_{n,i}$) and outputs ($y_{1,i}$, $y_{2,i}$, ..., $y_{m,i}$)?
- Least squares not tractable
- Data may not appear to be a function
 - Experimental variability, or measuring multiple effects?



Model fitting

- What if we have a data set of inputs ($x_{1,i}$, $x_{2,i}$, ..., $x_{n,i}$) and outputs ($y_{1,i}$, $y_{2,i}$, ..., $y_{m,i}$)?
- Least squares not tractable
- Data may not appear to be a function
 - Experimental variability, or measuring multiple effects?



Machine learning (ML)

- “A computer program is said to learn from experience E with respect to some class of tasks T and performance measure P , if its performance at tasks in T , as measured by P , improves with experience E .” – Mitchell (1997)
- An algorithm that learns from data, without prior knowledge about the data
 - ML models are blank slates
 - ML is not magic that solves all of our problems

Types of ML


- Supervised learning
 - Models labeled data sets of input-output pairs
- Unsupervised learning
 - Models unlabeled data sets of outputs
- Semi-supervised learning
 - Combination of labeled and unlabeled data
- Reinforcement learning
 - Seeks to maximize reward or minimize risk

Supervised ML

- Training data is used to teach the model what the output should be for a given input
 - The data is fed to the model in batches
 - Model makes predictions on each batch
 - Small adjustments to parameters are made each time based on the error of its output as measured by the loss function (backpropagation)
- Fits model parameters to approximate a function, similar to least squares

Unsupervised ML


- Finds patterns in an unlabeled data set (outputs with no corresponding inputs)
- Common uses: clustering, anomaly detection



"An Introduction to Statistical Learning" by James, Witten, Hastie, and Tibshirani

Reinforcement ML

- Agent seeks to accomplish some task
- For each action the agent takes, the interpreter evaluates the reward function
- Agent's future actions depend on rewards of past states




ML Task Examples

- Classification
- Regression (predict a value given some input)
- Transcription
- Translation




<https://www.shutterstock.com/image-photo/dog-cat-179786294>



Mnist database of handwritten digits

Neural networks (NNs)


- Type of ML model
- Uses sequential layers of nodes, where the outputs of nodes from one layer are the inputs for the next layer's nodes



Andrej Karpathy

Neural networks (NNs)


- Each node has an activation function to transform the inputs (model parameters)
- Each connection between nodes has an associated weight (model parameters)



Andrej Karpathy

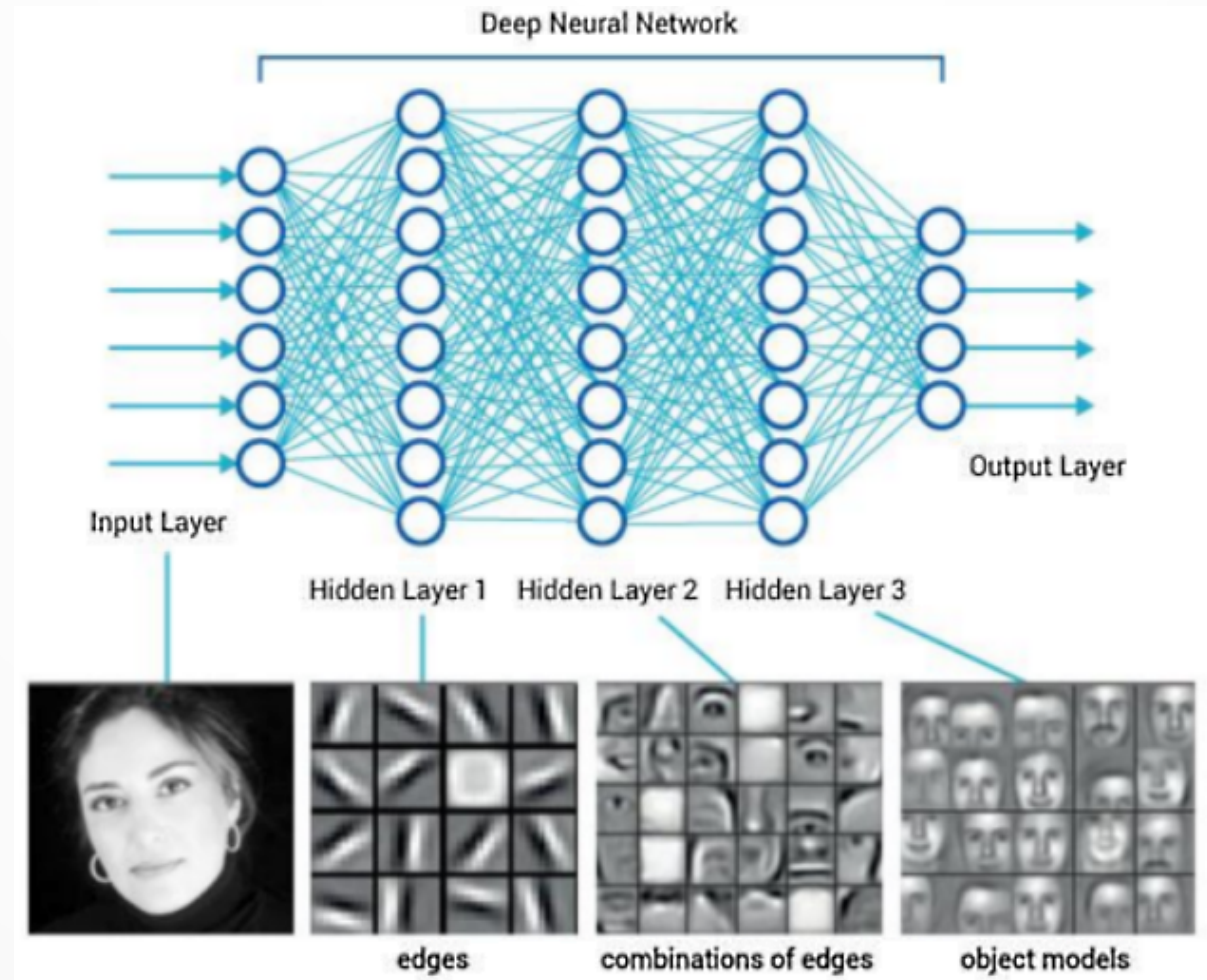
Neural networks (NNs)

- Nodes work similar to Boolean logic gates
- “Hidden layers” work on higher-order features in the data
- No hidden layers and linear activation functions?
Linear regression!
 - $\mathbf{y}_i = \mathbf{w} \cdot \mathbf{x}_i$
 - $\mathbf{y}_i = w_1x_1 + w_2x_2 + \dots$



Andrej Karpathy

What are hidden layers doing?



<https://iq.opengenus.org/hidden-layers/>


Activation Functions

- Controls how node inputs are mapped to outputs
 - Can be any function, as long as it is defined for $(-\infty, \infty)$
- Common activations:
 - Linear (identity)
 - Exponential
 - Softmax
 - Softplus
 - Scaled ELU

https://miro.medium.com/max/1192/1*4ZEDRpFuClpUjNgjDdT2Lg.png


Sigmoid

$$\sigma(x) = \frac{1}{1+e^{-x}}$$




tanh

$$\tanh(x)$$



ReLU


$$\max(0, x)$$



Rectified Linear Unit

Leaky ReLU

$$\max(0.1x, x)$$




Maxout

$$\max(w_1^T x + b_1, w_2^T x + b_2)$$

ELU

$$\begin{cases} x & x \geq 0 \\ \alpha(e^x - 1) & x < 0 \end{cases}$$




Exponential Linear Unit

NN Math Example

- NNs are just doing weird math... Let's see it
- Assume linear activation for **y**, sigmoid activation for **h₂**, and tanh activation for **h₁**

- $\mathbf{y} = \mathbf{w}_3 \cdot \mathbf{h}_2$
- $\mathbf{h}_2 = \{1 + \exp(-\mathbf{w}_2 \cdot \mathbf{h}_1)\}^{-1}$
- $\mathbf{h}_1 = \tanh(\mathbf{w}_1 \cdot \mathbf{x})$
- $\mathbf{y} = \mathbf{w}_3 \cdot \{1 + \exp(-\mathbf{w}_2 \cdot \tanh(\mathbf{w}_1 \cdot \mathbf{x}))\}^{-1}$



Andrej Karpathy

Demo: visual NN

- Before we get into more details...
- <https://playground.tensorflow.org>

How to train an NN

- A data set is split into training, validation, and test sets
 - Training set: trains the network and updates the weights
 - Validation set: monitors training and prevents overfitting
 - Test set: evaluates model performance/generalization

How to train an NN


- Loss functions govern how an NN learns to minimize the error in predictions
- Common losses:
 - Mean squared error
 - Mean absolute error
 - Binary crossentropy (for binary classification)
 - Categorical crossentropy (for multiclassification)
- You can also define your own!

How to train an NN

- Training steps
 - Make a prediction for some input
 - Calculate the error according to loss function
 - Update the weights via backpropagation
- Validation steps
 - Tests the generality of the model during training
 - Halts training when overfitting detected (early stopping)


How to train an NN

- Every N training steps, make 1 validation step
- Repeat for many epochs (iterations through the training set) to learn everything that it can, without overfitting
- Then, the model can be applied to the test set, which evaluates the performance of the model



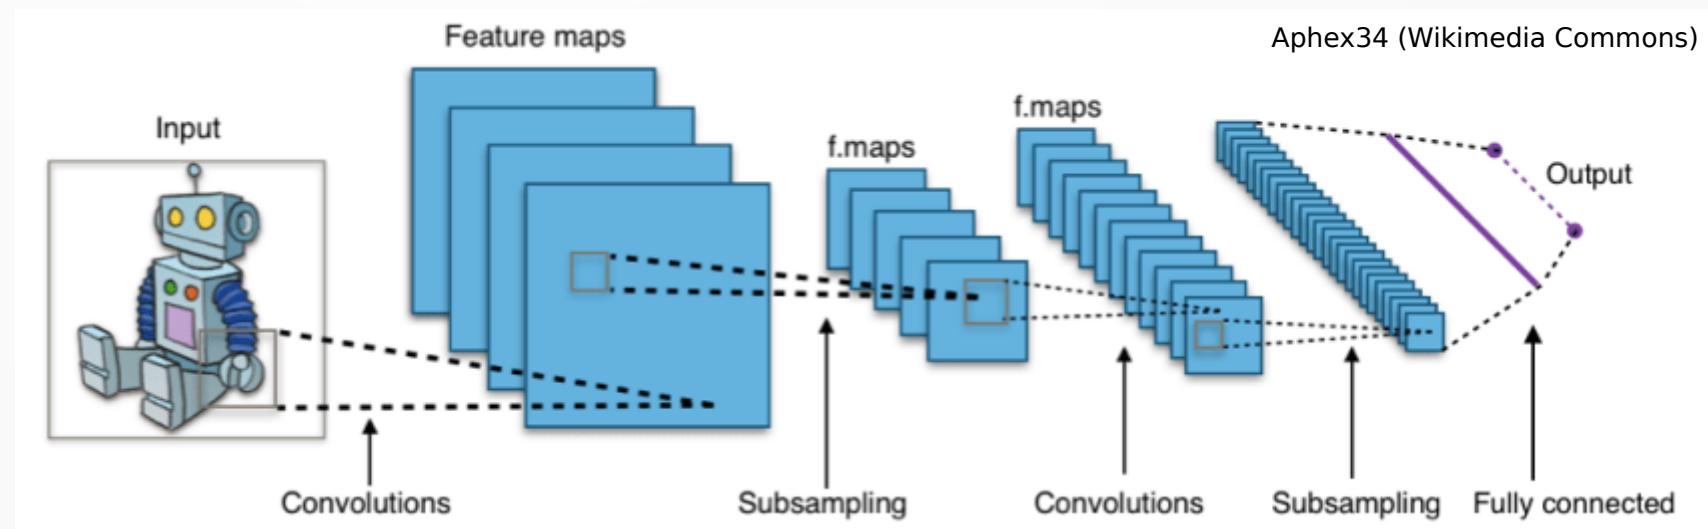

Feedforward/Dense NNs

- Simplest NN, where nodes are connected to all nodes in adjacent layers
- Useful in tasks like image classification and regression
- In theory, they can model anything!
 - In practice, other approaches are better for certain tasks



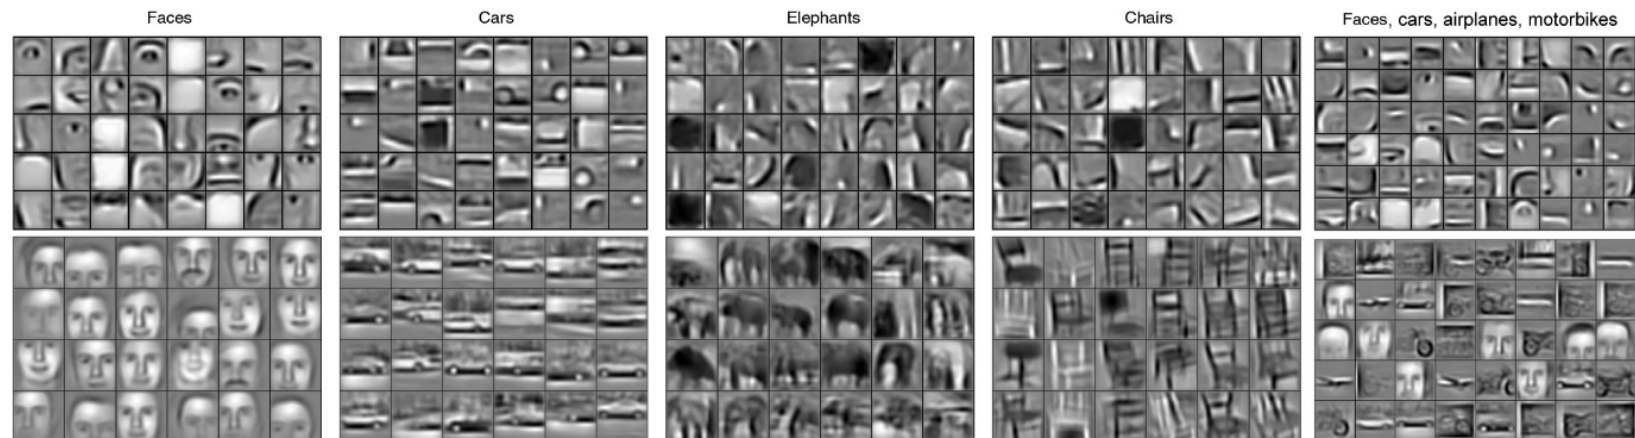

Convolutional NNs

- Nodes are connected to only a few nodes in adjacent layers
 - Like collection of dense NNs, where each NN is highly specialized (feature maps)
- Useful in classification, regression, computer vision, signal/image processing



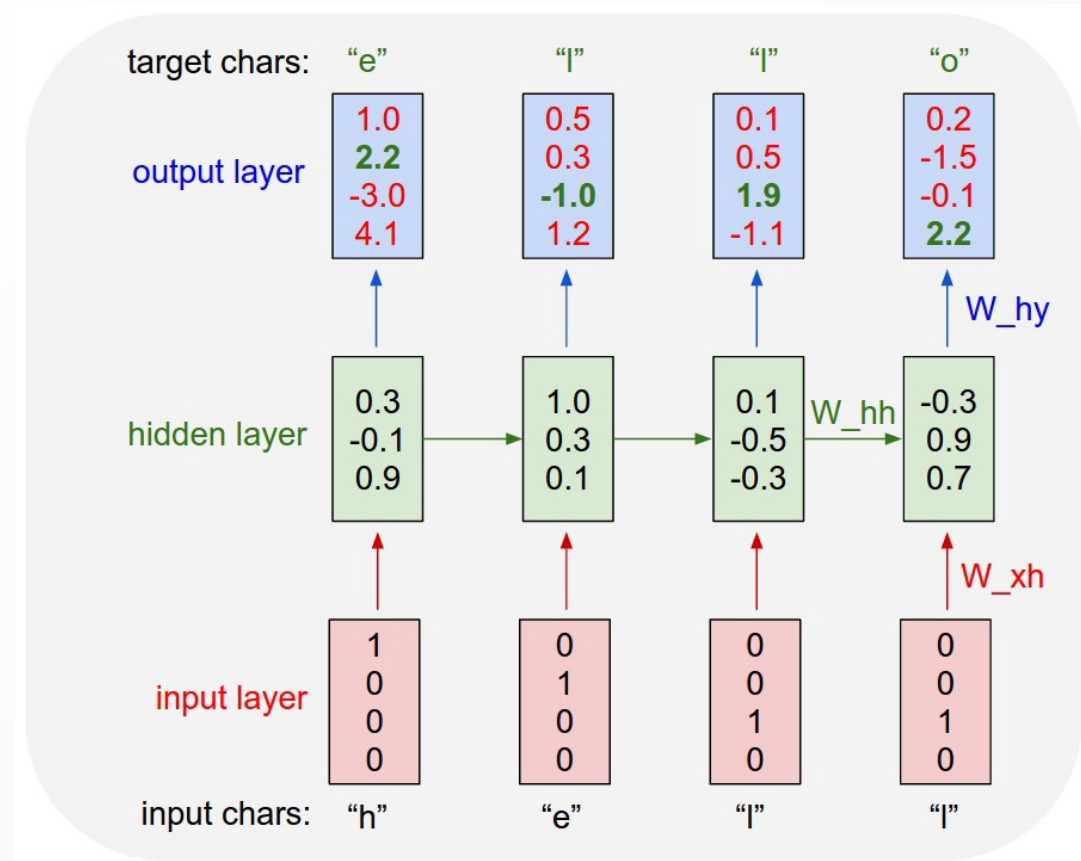
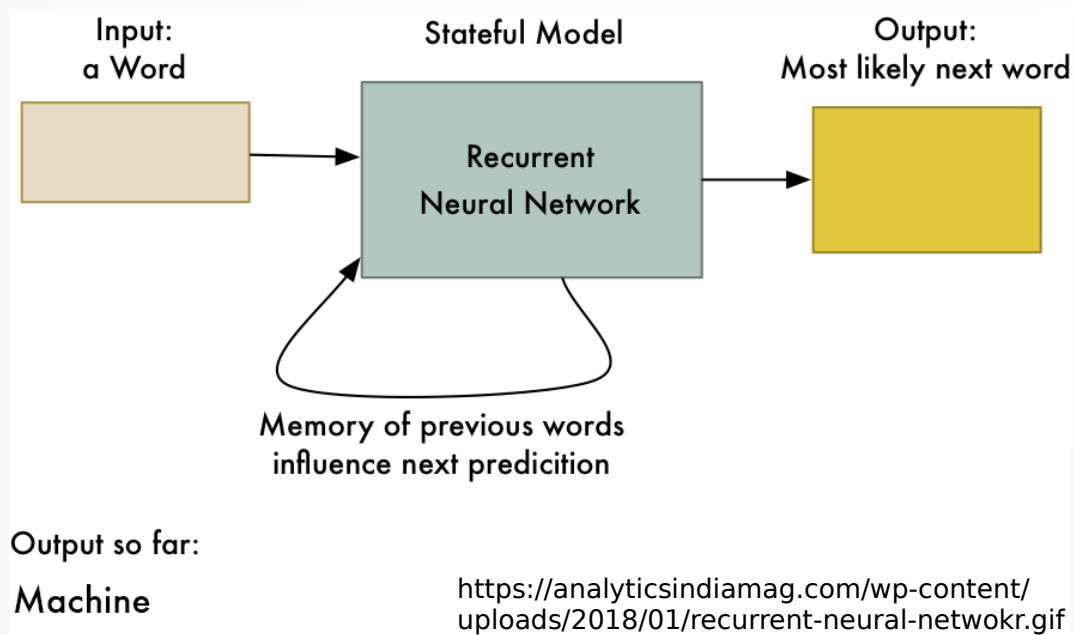
Convolutional NNs

- Nodes are connected to only a few nodes in adjacent layers
 - Like collection of dense NNs, where each NN is highly specialized (feature maps)
- Useful in classification, regression, computer vision, signal/image processing



Recurrent NNs

- Stores information about previous inputs to inform future predictions
- Useful for time-series data



https://miro.medium.com/max/902/1*IMalbwl6uj3nlqxixZYFvA.jpeg

Applications

- Automated driving gets a lot of attention
- But, there are a lot of other useful applications...
- <https://www.youtube.com/watch?v=D5VN56jQMWM>
- <https://www.youtube.com/watch?v=gn4nRCC9TwQ>
- <https://www.youtube.com/watch?v=fUyU3IKzoio>
- <https://www.youtube.com/watch?v=aFuA50H9uek>
- If interested, some more applications:
- <https://www.youtube.com/watch?v=HKcO3-6TYr0>
- <https://www.youtube.com/watch?v=Bui3DWs02h4>

Further resources

- Online
 - Image classification: <https://distill.pub/2018/building-blocks/>
 - Build NN with Numpy:
<https://towardsdatascience.com/neural-net-from-scratch-using-numpy-71a31f6e3675>
- Books
 - *Deep Learning* by Goodfellow, Bengio, and Courville
 - *Pattern Recognition and Machine Learning* by Bishop
- Internet searches! There are many great tutorials available

Demo: basic dense NNs

- Model a unit circle with parameter θ
- Model a Gaussian with parameters position, mean, and standard deviation
- Of course, neither of these actually need ML to solve – we're just demonstrating how to use ML
 - Most problems require graphics processing units (GPUs) in order to train in a reasonable amount of time